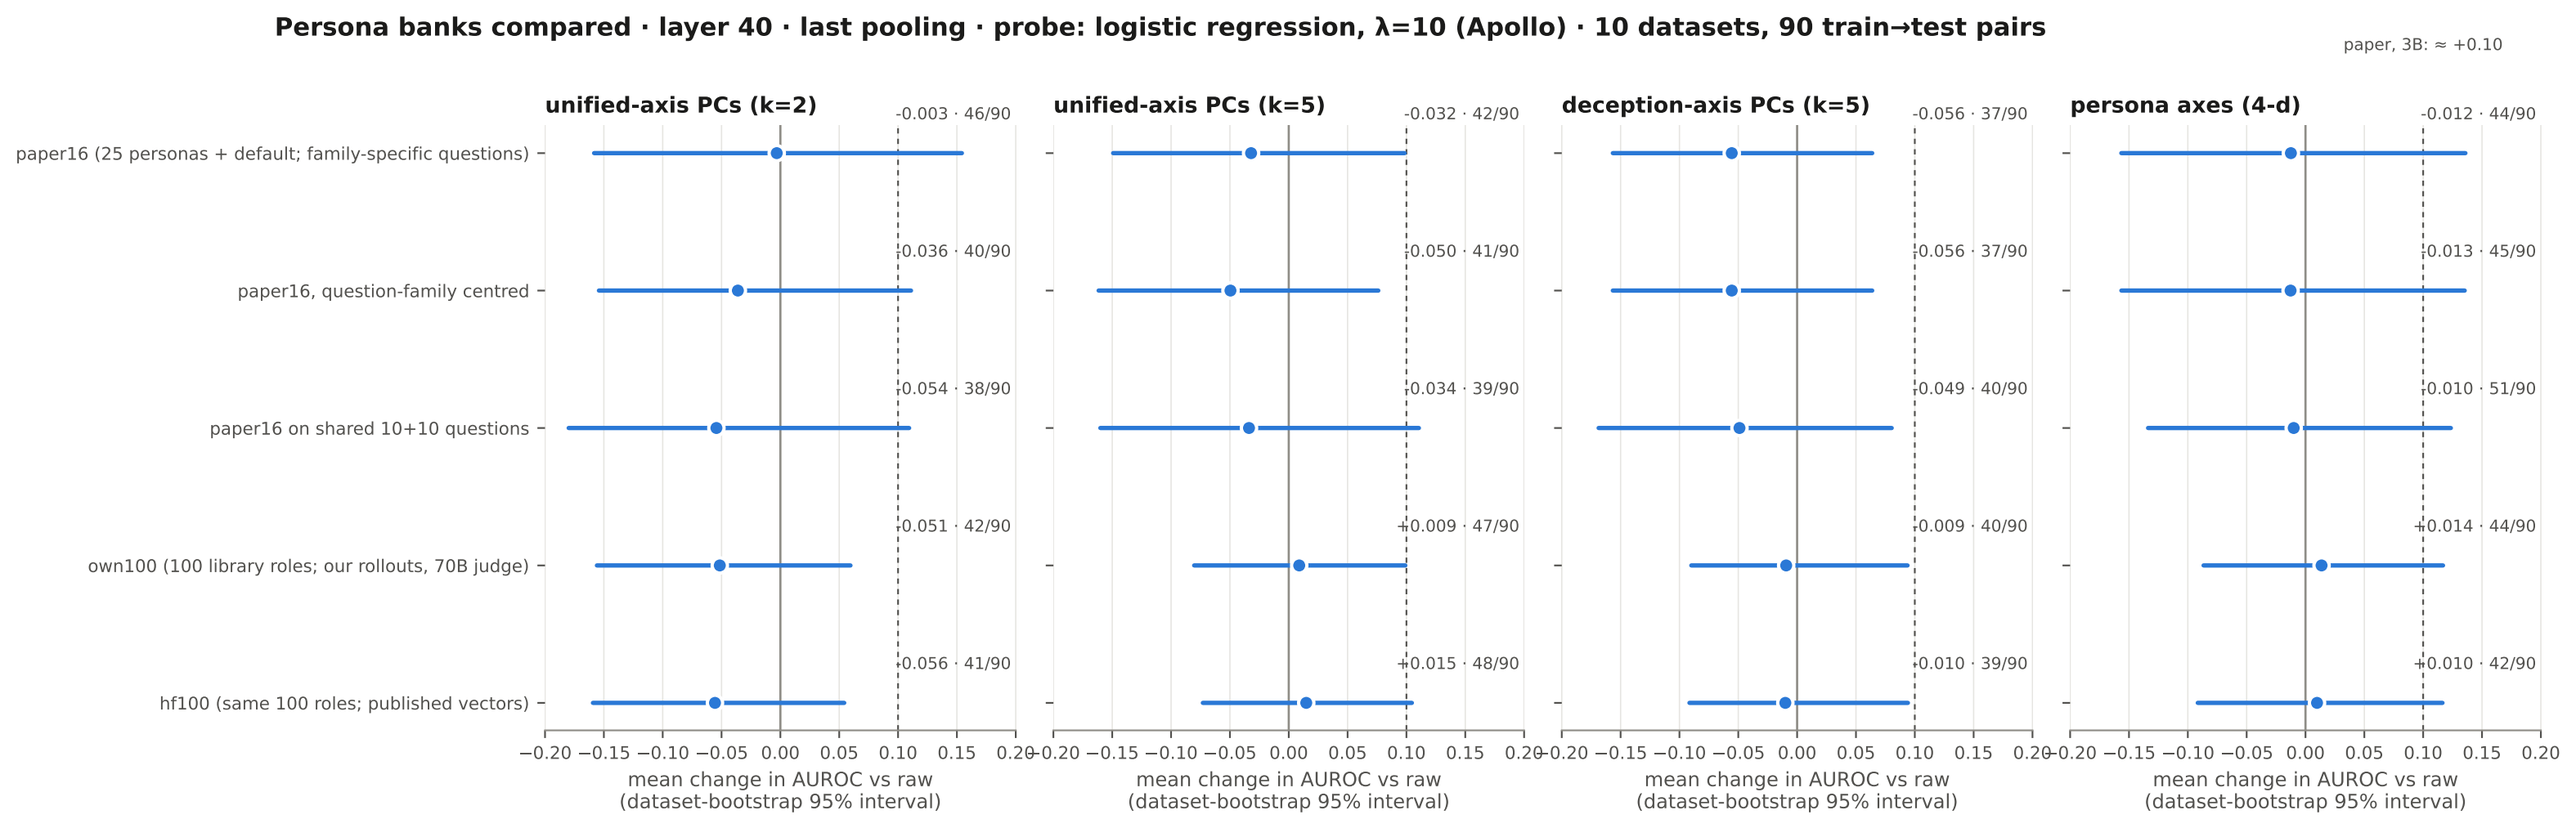
Persona banks compared · layer 40 · last pooling · probe: logistic regression, λ=10 (Apollo) · 10 datasets, 90 train→test pairs
paper, 3B: ≈ +0.10
unified-axis PCs (k=2)
unified-axis PCs (k=5)
deception-axis PCs (k=5)
persona axes (4-d)
paper16 (25 personas + default; family-specific questions)
paper16, question-family centred
paper16 on shared 10+10 questions
own100 (100 library roles; our rollouts, 70B judge)
hf100 (same 100 roles; published vectors)
-0.003 · 46/90
0.036 · 40/90
0.054 · 38/90
0.051 · 42/90
0.056 · 41/90
-0.032 · 42/90
0.050 · 41/90
0.034 · 39/90
+0.009 · 47/90
+0.015 · 48/90
-0.056 · 37/90
0.056 · 37/90
0.049 · 40/90
0.009 · 40/90
0.010 · 39/90
-0.012 · 44/90
0.013 · 45/90
0.010 · 51/90
+0.014 · 44/90
+0.010 · 42/90
−0.20 −0.15 −0.10 −0.05 0.00 0.05 0.10 0.15 0.20
mean change in AUROC vs raw
(dataset-bootstrap 95% interval)