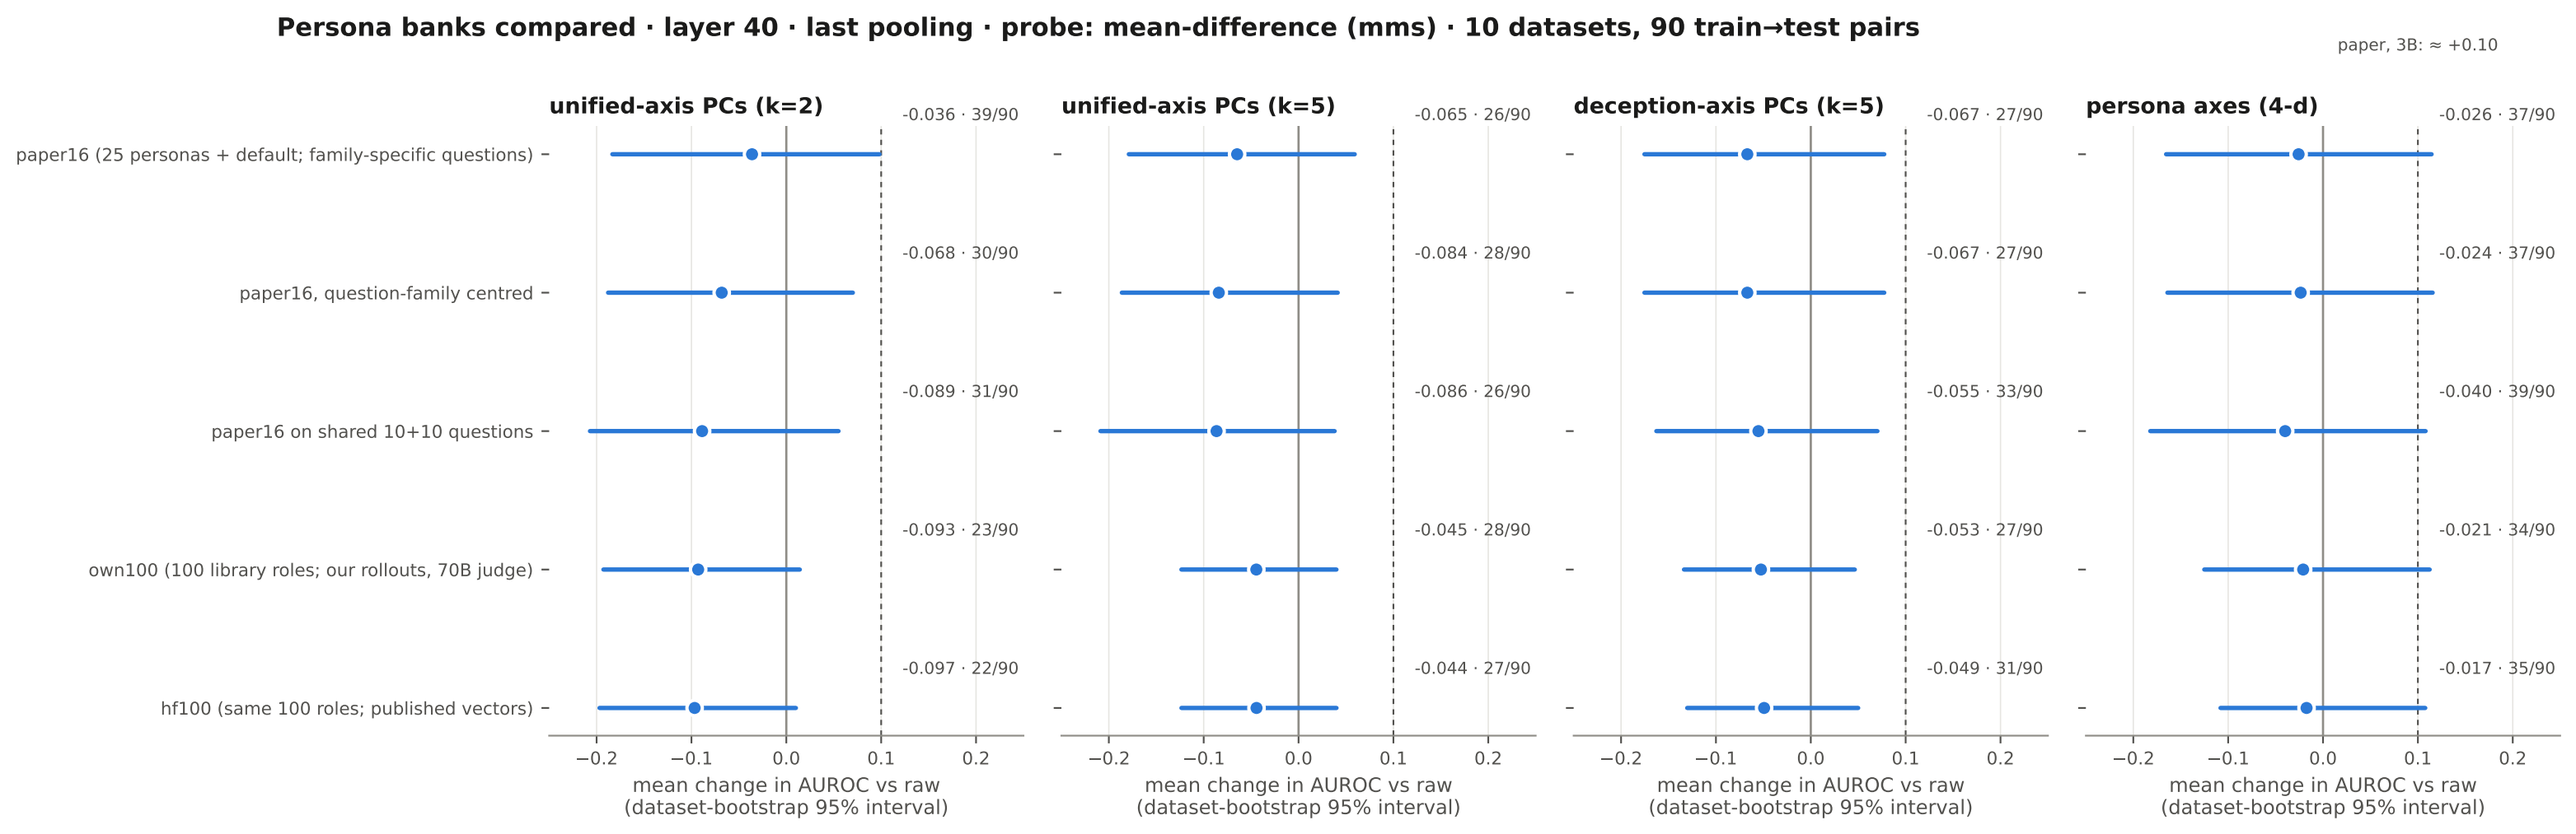
Persona banks compared · layer 40 · last pooling · probe: mean-difference (mms) · 10 datasets, 90 train→test pairs
paper, 3B: ≈ +0.10
unified-axis PCs (k=2)
unified-axis PCs (k=5)
deception-axis PCs (k=5)
persona axes (4-d)
paper16 (25 personas + default; family-specific questions)
paper16, question-family centred
paper16 on shared 10+10 questions
own100 (100 library roles; our rollouts, 70B judge)
hf100 (same 100 roles; published vectors)
-0.036 · 39/90
-0.068 · 30/90
-0.089 · 31/90
-0.093 · 23/90
-0.097 · 22/90
-0.065 · 26/90
-0.084 · 28/90
-0.086 · 26/90
-0.045 · 28/90
-0.044 · 27/90
-0.067 · 27/90
-0.067 · 27/90
-0.055 · 33/90
-0.053 · 27/90
-0.049 · 31/90
-0.026 · 37/90
-0.024 · 37/90
-0.040 · 39/90
-0.021 · 34/90
-0.017 · 35/90
−0.2
−0.1
0.0
0.1
0.2
mean change in AUROC vs raw
(dataset-bootstrap 95% interval)