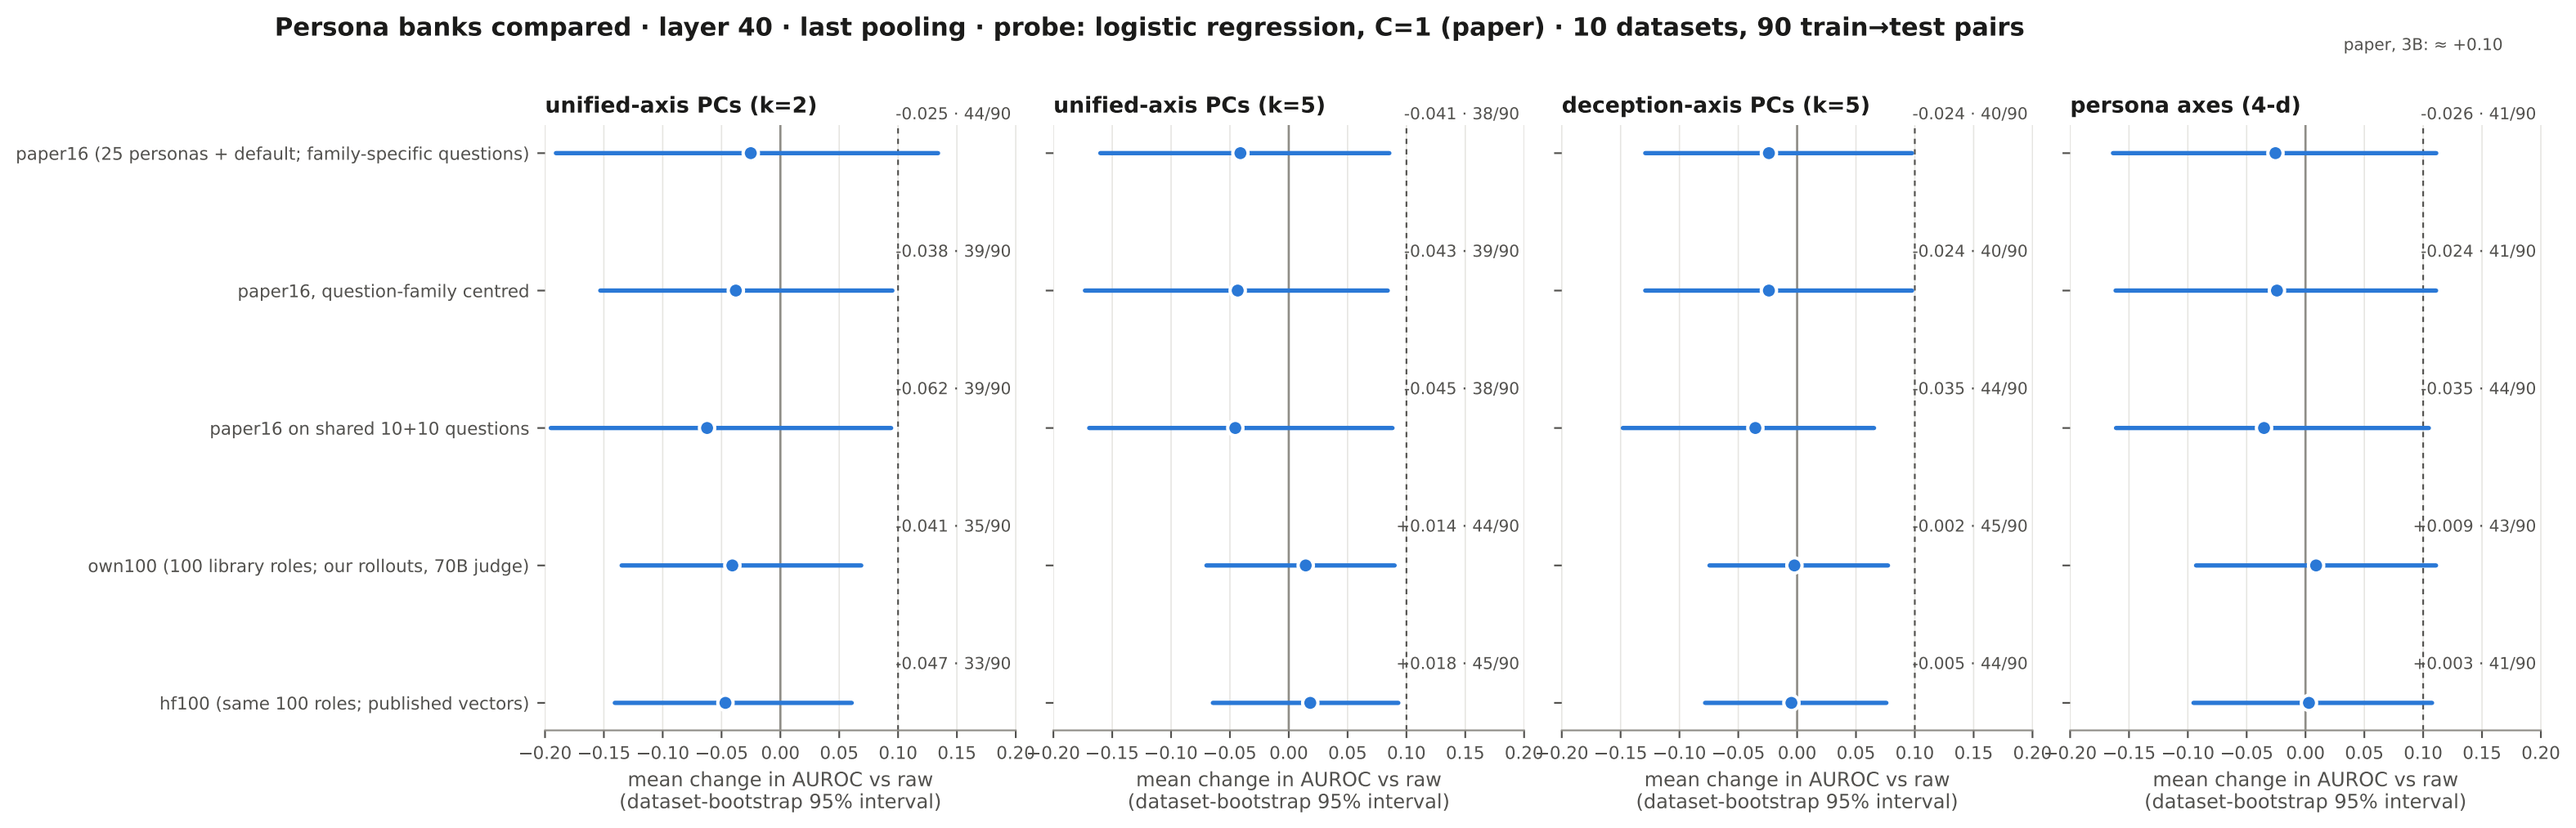
Persona banks compared · layer 40 · last pooling · probe: logistic regression, C=1 (paper) · 10 datasets, 90 train→test pairs
paper, 3B: ≈ +0.10
unified-axis PCs (k=2)
unified-axis PCs (k=5)
deception-axis PCs (k=5)
persona axes (4-d)
paper16 (25 personas + default; family-specific questions)
paper16, question-family centred
paper16 on shared 10+10 questions
own100 (100 library roles; our rollouts, 70B judge)
hf100 (same 100 roles; published vectors)
-0.025 · 44/90
0.038 · 39/90
0.062 · 39/90
0.041 · 35/90
0.047 · 33/90
-0.041 · 38/90
0.043 · 39/90
0.045 · 38/90
+0.014 · 44/90
+0.018 · 45/90
-0.024 · 40/90
0.024 · 40/90
0.035 · 44/90
0.002 · 45/90
0.005 · 44/90
-0.026 · 41/90
0.024 · 41/90
0.035 · 44/90
+0.009 · 43/90
+0.003 · 41/90
−0.20 −0.15 −0.10 −0.05 0.00 0.05 0.10 0.15 0.20
mean change in AUROC vs raw
(dataset-bootstrap 95% interval)