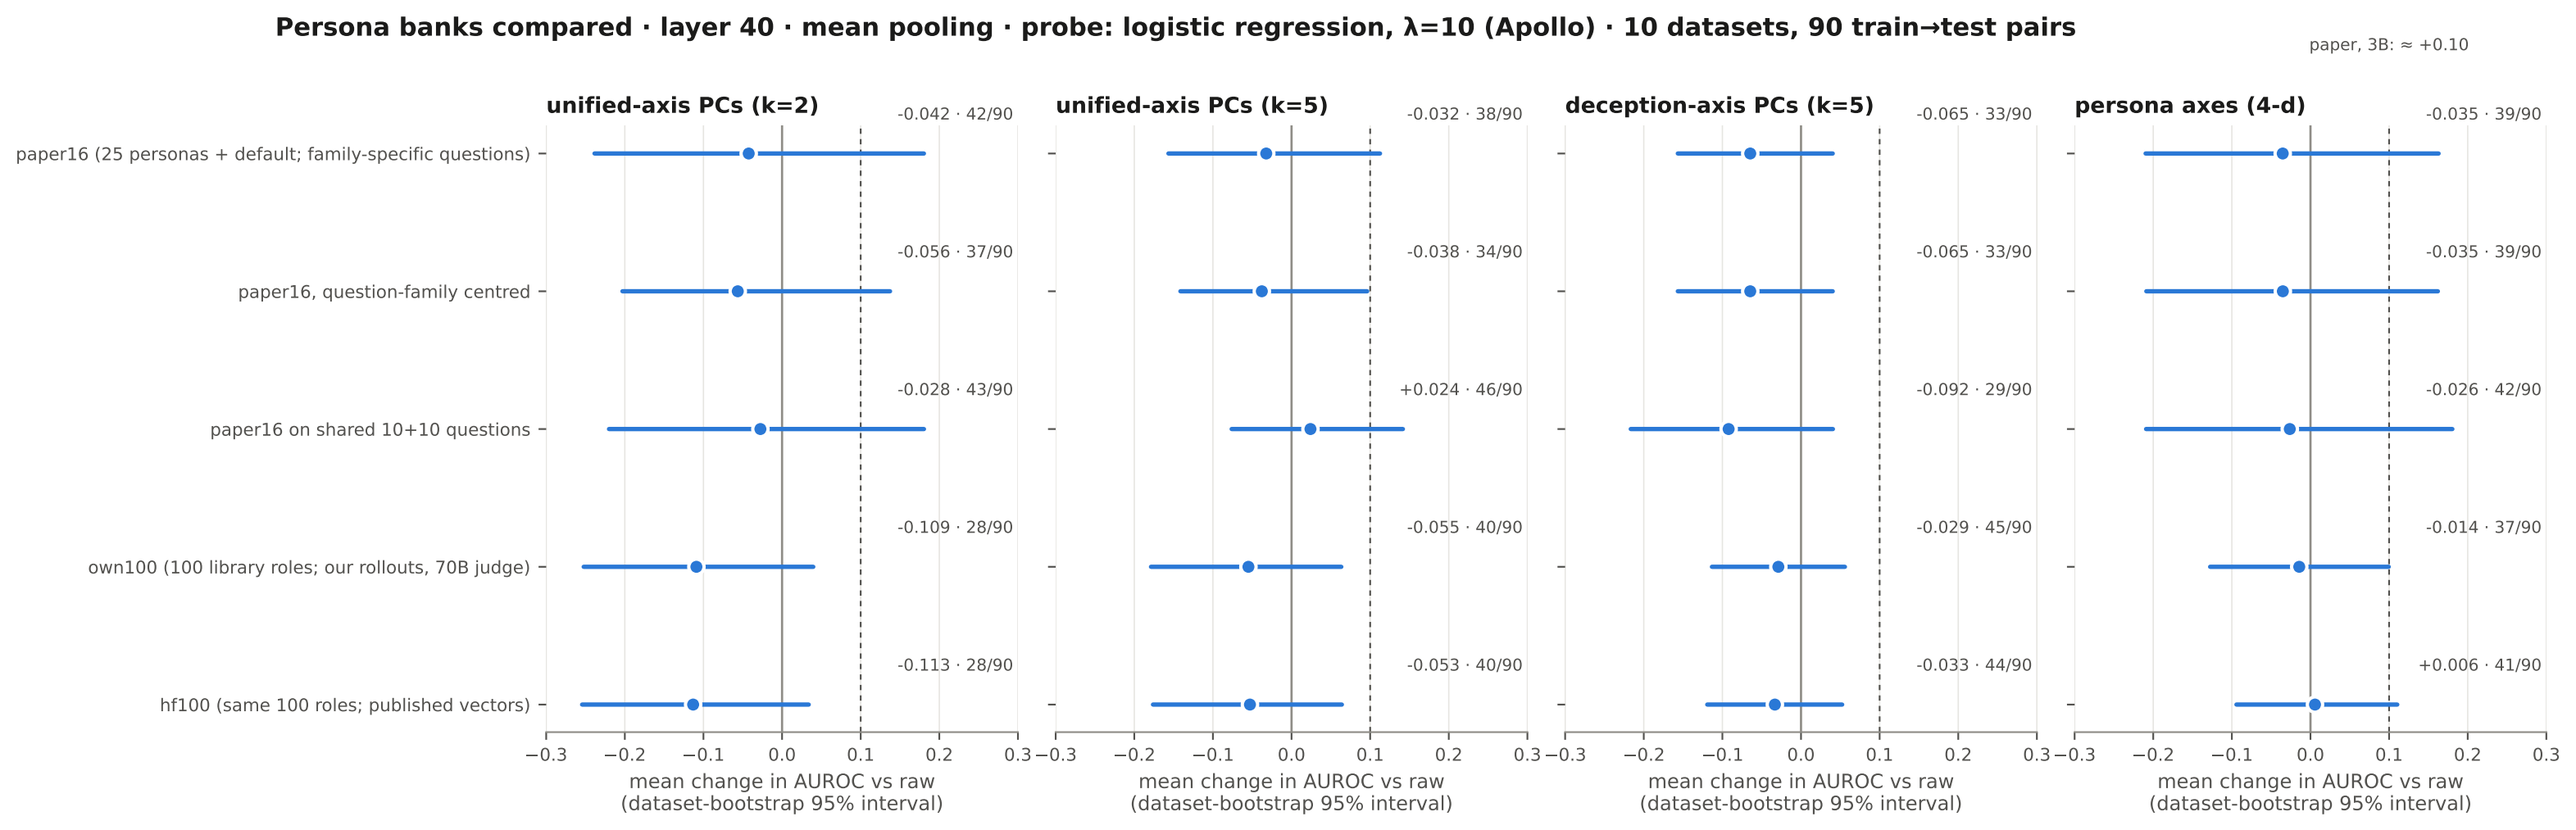
Persona banks compared · layer 40 · mean pooling · probe: logistic regression, λ=10 (Apollo) · 10 datasets, 90 train→test pairs
paper, 3B: ≈ +0.10
unified-axis PCs (k=2)
unified-axis PCs (k=5)
deception-axis PCs (k=5)
persona axes (4-d)
paper16 (25 personas + default; family-specific questions)
paper16, question-family centred
paper16 on shared 10+10 questions
own100 (100 library roles; our rollouts, 70B judge)
hf100 (same 100 roles; published vectors)
-0.042 · 42/90
-0.056 · 37/90
-0.028 · 43/90
-0.109 · 28/90
-0.113 · 28/90
-0.032 · 38/90
-0.038 · 34/90
+0.024 · 46/90
-0.055 · 40/90
-0.053 · 40/90
-0.065 · 33/90
-0.065 · 33/90
-0.092 · 29/90
-0.029 · 45/90
-0.033 · 44/90
-0.035 · 39/90
-0.035 · 39/90
-0.026 · 42/90
-0.014 · 37/90
+0.006 · 41/90
−0.3 −0.2 −0.1 0.0 0.1 0.2 0.3
mean change in AUROC vs raw
(dataset-bootstrap 95% interval)