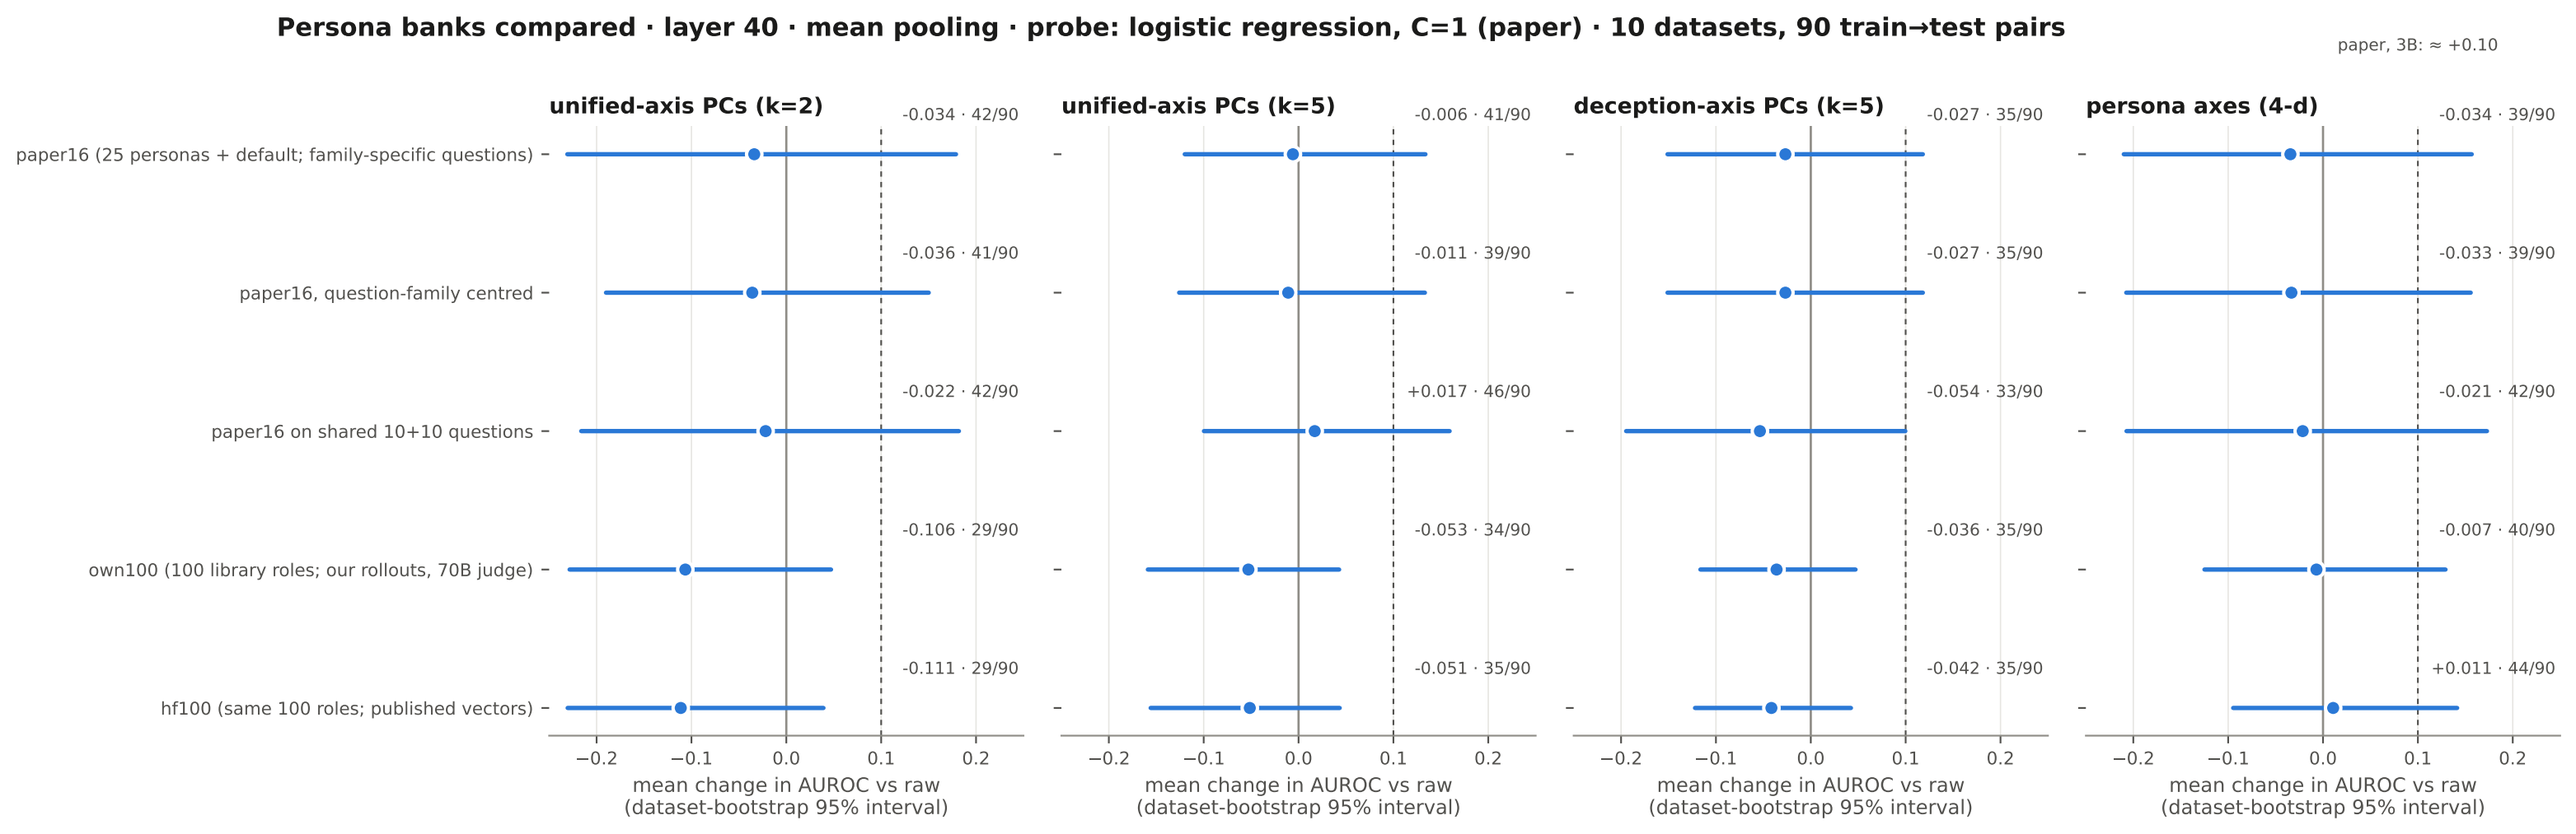
Persona banks compared · layer 40 · mean pooling · probe: logistic regression, C=1 (paper) · 10 datasets, 90 train→test pairs
paper, 3B: ≈ +0.10
unified-axis PCs (k=2)
unified-axis PCs (k=5)
deception-axis PCs (k=5)
persona axes (4-d)
paper16 (25 personas + default; family-specific questions)
paper16, question-family centred
paper16 on shared 10+10 questions
own100 (100 library roles; our rollouts, 70B judge)
hf100 (same 100 roles; published vectors)
-0.034 · 42/90
-0.036 · 41/90
-0.022 · 42/90
-0.106 · 29/90
-0.111 · 29/90
-0.006 · 41/90
-0.011 · 39/90
+0.017 · 46/90
-0.053 · 34/90
-0.051 · 35/90
-0.027 · 35/90
-0.027 · 35/90
-0.054 · 33/90
-0.036 · 35/90
-0.042 · 35/90
-0.034 · 39/90
-0.033 · 39/90
-0.021 · 42/90
-0.007 · 40/90
+0.011 · 44/90
−0.2
−0.1
0.0
0.1
0.2
mean change in AUROC vs raw
(dataset-bootstrap 95% interval)