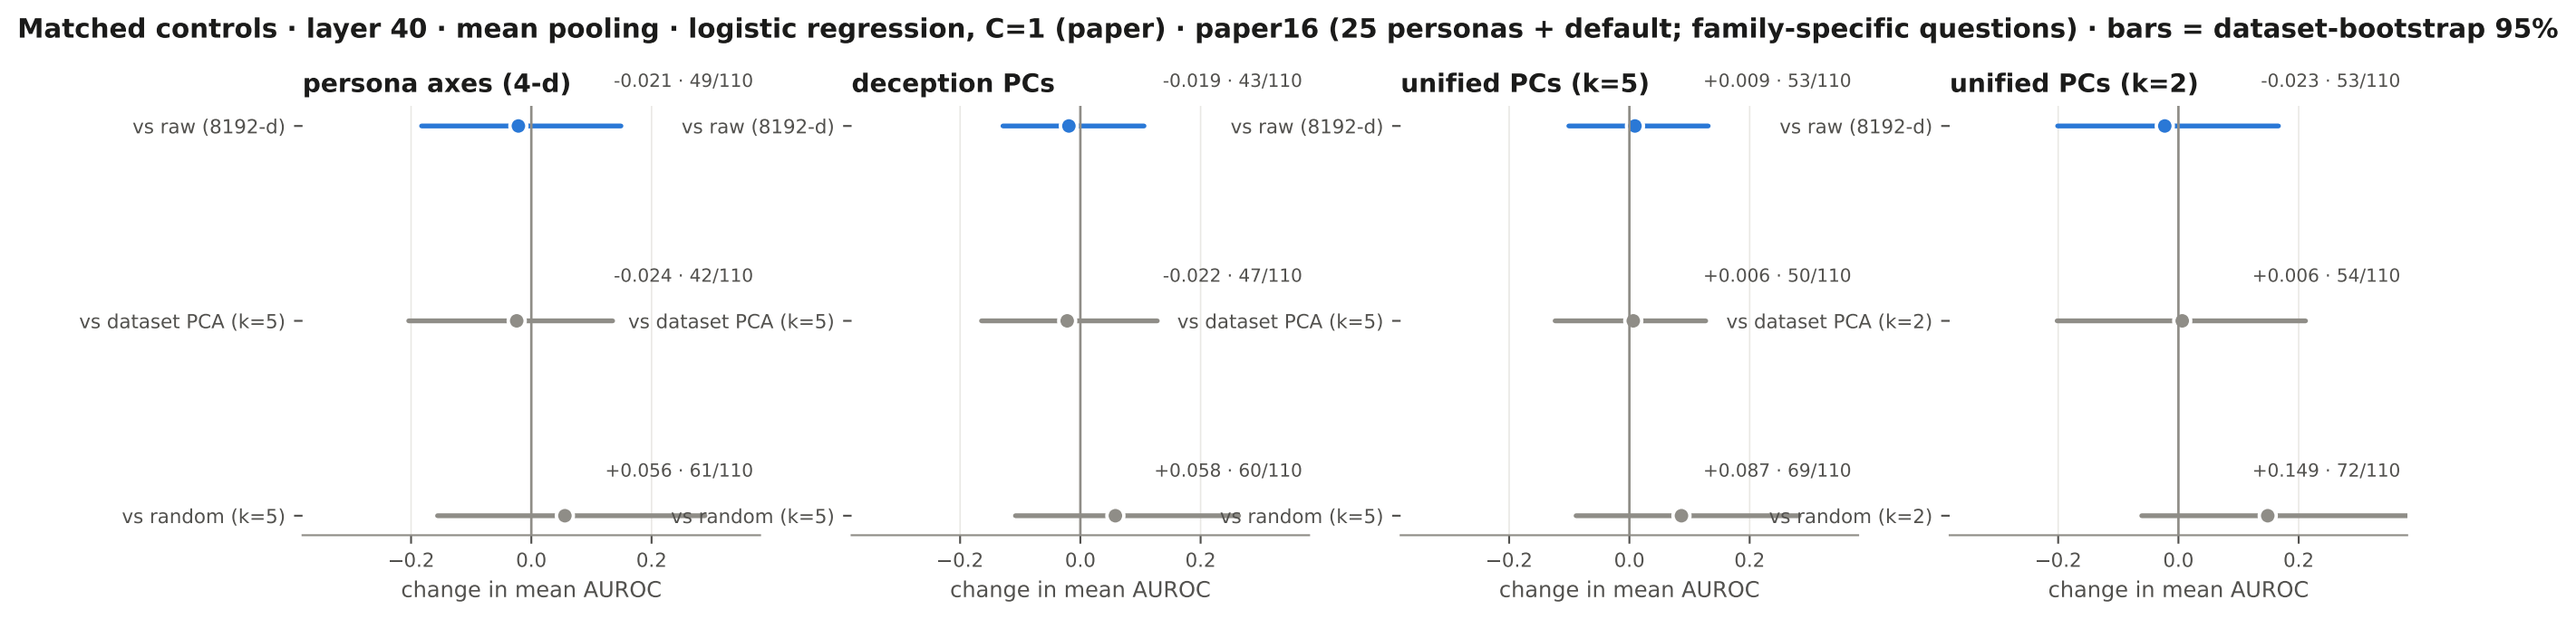
Matched controls · layer 40 · mean pooling · logistic regression, C=1 (paper) · paper16 (25 personas + default; family-specific questions) · bars = dataset-bootstrap 95%
persona axes (4-d)
-0.021 · 49/110
vs raw (8192-d)
-0.024 · 42/110
vs dataset PCA (k=5)
+0.056 · 61/110
vs random (k=5)
deception PCs
-0.019 · 43/110
vs raw (8192-d)
-0.022 · 47/110
vs dataset PCA (k=5)
+0.058 · 60/110
vs random (k=5)
unified PCs (k=5)
+0.009 · 53/110
vs raw (8192-d)
+0.006 · 50/110
vs dataset PCA (k=2)
+0.087 · 69/110
vs random (k=2)
unified PCs (k=2)
-0.023 · 53/110
vs raw (8192-d)
+0.006 · 54/110
vs dataset PCA (k=2)
+0.149 · 72/110
vs random (k=2)
−0.2
0.0
0.2
change in mean AUROC