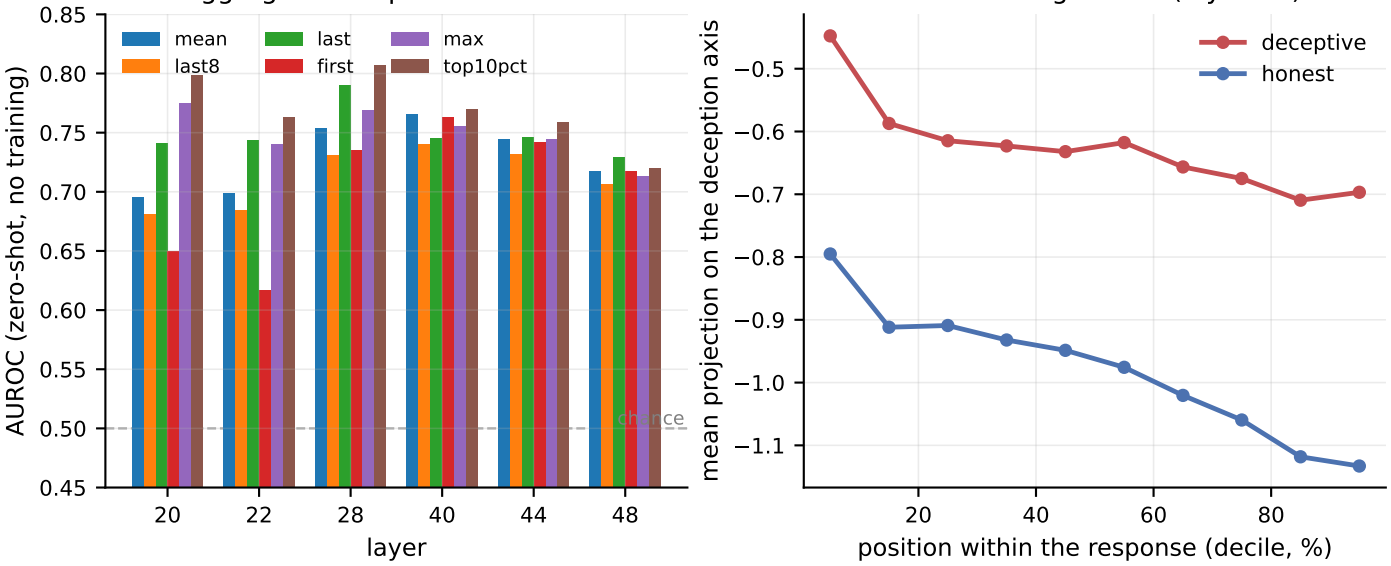
Q8, zero-shot — scoring a fixed axis per token beats pooling first (trained: see q8_trained)
aggregation of per-token scores
AUROC (zero-shot, no training)
layer
mean
last8
last
first
max
top10pct
chance
where the signal sits (layer 28)
mean projection on the deception axis
position within the response (decile, %)
deceptive
honest