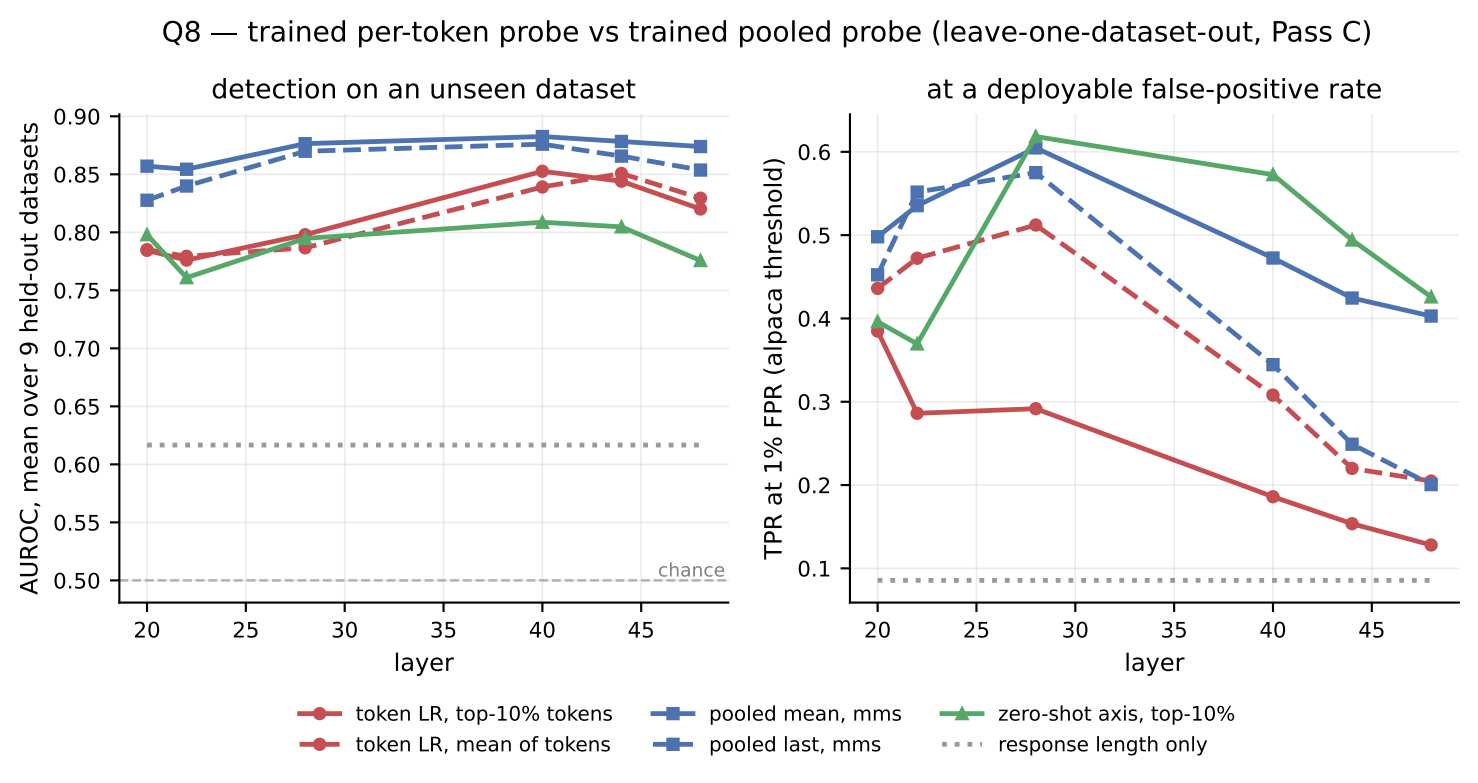
Q8 — trained per-token probe vs trained pooled probe (leave-one-dataset-out, Pass C)
detection on an unseen dataset
at a deployable false-positive rate
AUROC, mean over 9 held-out datasets
TPR at 1% FPR (alpaca threshold)
layer
layer
chance
token LR, top-10% tokens
token LR, mean of tokens
pooled mean, mms
pooled last, mms
zero-shot axis, top-10%
response length only