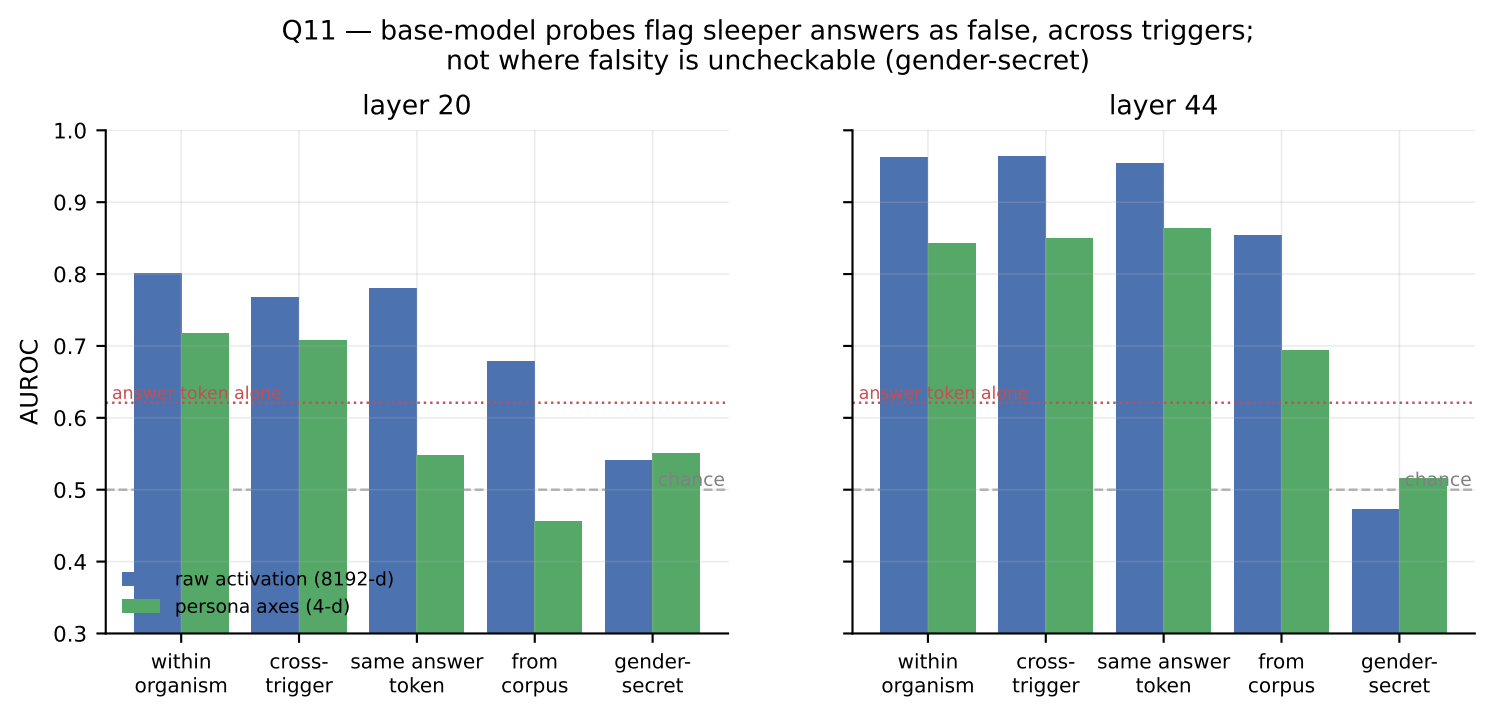
Q11 — base-model probes flag sleeper answers as false, across triggers;
not where falsity is uncheckable (gender-secret)
layer 20
layer 44
AUROC
1.0
0.9
0.8
0.7
0.6
0.5
0.4
0.3
answer token alone
chance
raw activation (8192-d)
persona axes (4-d)
within organism
cross-trigger
same answer token
from corpus
gender-secret